

SLM Guard

Guardrails for Specialist Small Models

Product Brief — v4 | August 2026

Revision note: v4 renames the product from SLM OS to SLM Guard. No positioning, capability, or roadmap content changed — this is a naming update only. All prior-version content changes (v2, v3) still apply and are carried forward under the new name.

The Problem

Specialist small language models deliver strong cost and latency advantages on high-volume banking workflows (fraud triage, compliance checking, alert review, dispute handling). In practice, three problems limit their adoption and long-term value:

1. **Drift** — Models degrade as fraud patterns, products, and policies change.
2. **Control risk** — Giving small models direct decision or tool access creates audit and operational risk.
3. **Data movement barriers** — Improving models from real production behavior usually requires moving sensitive data, which is blocked or heavily restricted.

Most institutions either keep expensive general models in the loop or accept that specialist models will slowly lose accuracy.

What SLM Guard Is

SLM Guard is the operating system for specialist small models in regulated decision workflows.

It enforces three non-negotiable properties:

- **Specialist models only recommend.** They never take action or call systems directly. A deterministic control plane validates every recommendation, applies policy, decides, and logs the outcome.
- **Continuous improvement without moving raw data.** Production signals are converted into privacy-safe training data so models can be specialized over time inside the controlled boundary.
- **Regulated promotion and full lineage.** New model versions are evaluated against explicit gates and only promoted when they meet accuracy, safety, and policy criteria. Every decision and every model version is fully traceable.

What the Control Plane Owns

This is stated explicitly because it is the core institutional-control argument for risk and audit stakeholders: the model recommends; it does not decide. The control plane — not the model — holds:

- **Final decision authority.** Every action taken is a control-plane decision, never a direct model action.
- **All tool and system access.** The model has no ability to call downstream systems, execute transactions, or trigger workflows.

- **Every resulting state change.** Nothing changes in a production system except through the control plane.
- **The complete audit record.** Every recommendation, policy check, decision, and outcome is logged there, not reconstructed after the fact.

How the Specialization Loop Works

This is the mechanism behind the "improve without moving data" claim, stated explicitly:

1. **Trace capture.** Every recommendation the specialist model makes — along with its confidence score, the policy checks applied, and the eventual outcome — is logged inside the institution's own boundary. Raw case data never leaves that boundary.
2. **Synthetic conversion.** Traces are converted into synthetic training examples that preserve the statistical and structural pattern of difficult, missed, or edge-case decisions without carrying the underlying customer or transaction data itself. Only these derived, privacy-safe examples are used for training.
3. **Light fine-tuning.** A lightweight adaptation (QLoRA-style) is run on the current specialist model using the synthetic pool, entirely inside the controlled environment.
4. **Held-out evaluation.** The candidate model version is scored against a frozen held-out test set it has never seen, plus a regression comparison against the current production version.
5. **Promotion gate.** The candidate is promoted only if it clears the criteria below. If it doesn't — or if there isn't yet enough new signal to justify a change — the current production version keeps running. No promotion is a normal, expected outcome, not a failure of the loop.

This loop is what lets accuracy improve over time without the raw-data-movement problem that blocks most model-improvement efforts in regulated environments.

Current Limits of the Privacy-Preserving Step

Stated plainly, because this claim will be the one a serious technical or risk reviewer presses hardest: in Phase 1, "privacy-safe" means careful synthetic rewriting under a defined quality process plus human review — it is not yet a formal mathematical privacy guarantee (e.g., differential privacy). This is an honest, defensible starting posture, not a finished one. Stronger, formally verifiable guarantees — and independent leakage testing of the conversion step — are planned as a hardening item in a later phase, once the Phase 1 loop is proven end to end. This brief will be updated when that work exists; it is not being claimed early.

Evaluation & Promotion Gate Criteria

A candidate model version must clear the following before it can replace the production version:

Gate	Criterion
Accuracy	No regression against the frozen held-out test set relative to current production baseline

Policy / safety compliance	Zero tolerance on recommendations that violate defined policy constraints
Confidence calibration	Tracked from cycle one; becomes a hard gate once a stable calibration baseline exists. In early cycles it is monitored, not blocking, since small-model confidence scores are frequently under-calibrated and a premature hard gate would either stall promotion indefinitely or mask the real signal
Regression / challenge set	No new failures introduced on the curated set of edge cases and prior known failure patterns

Every evaluation run, pass/fail result, and promotion decision — including a "no promotion this cycle" outcome — is logged and replayable.

Core Capabilities

Capability	Enterprise Value
Structured recommendation contract	Models cannot bypass controls or invent actions
Confidence-based escalation	Low-confidence or high-risk cases route to larger models or humans; most volume stays cheap
Deterministic control plane	Policy, decision rights, and audit remain with the institution
Privacy-preserving specialization loop	Models improve from real usage without raw data leaving the boundary
Evaluation & promotion gates	Model risk management and compliance get clear, inspectable criteria
Complete decision lineage	Every recommendation, confidence score, policy check, and outcome is replayable

Primary Use Cases (Initial Focus)

- Fraud alert triage and prioritization
- Compliance policy checking and exception handling
- High-volume operational decision support where accuracy, cost, and auditability must coexist

Who It Is For

- Banks and large credit unions running high-volume, policy-sensitive workflows
- AI, Fraud, Compliance, and Operations leaders who need specialist models to stay accurate without increasing risk or cost
- Model Risk Management and Audit stakeholders who require clear lineage and controlled promotion

What Makes It Different

Most current approaches to specialist small models stop at one of two points:

- **One-time domain fine-tuning**, delivered as a tuned model with no built-in mechanism to keep it current as patterns shift, and no privacy-safe path to retrain it from real usage.
- **Generic model hosting / MLOps platforms**, which manage deployment and monitoring but leave the questions of decision authority, privacy-safe improvement, and regulated promotion to the institution to solve separately.

SLM Guard is different because it owns all three as one integrated system: constrained authority (the model can never act outside its recommendation contract), privacy-preserving continuous improvement (the model gets better without raw data ever leaving the boundary), and regulated promotion (every version change is gated and every decision is auditable). No single piece of this is unprecedented in isolation — the differentiation is that it is delivered as one operating layer instead of three separate problems the institution has to integrate itself.

Deployment Posture

- Supports quantized specialist models running in private / on-premises environments
- Control plane and audit logs remain under institutional control
- Designed for staged adoption: start with one workflow, then expand

Current Status & Roadmap

Where this stands today: SLM Guard is in Phase 1 — building the first complete vertical reference implementation on Fraud Alert Triage. No production pilot has run yet; the claims in this brief describe what the system is designed to do, not yet a measured result.

Phase	Focus	Exit criterion
1 — Vertical reference implementation	Fraud Alert Triage: schema, control plane, baseline specialist model, first specialization loop	Complete loop runs end-to-end with a measured before/after accuracy delta and full audit trail
2 — Harden into platform core	Extract control plane, specialization engine, and evaluation harness into reusable, versioned components; add model registry	Fraud Triage runs as a "domain pack" on top of a reusable core
3 — Platform shape & credibility	Open-source the core + reference implementation; documentation and architecture decision records	Public, reproducible reference deployment
4 — Enterprise packaging	Open-core, implementation accelerator, or internal offering model	Defined commercial engagement model

What Phase 1 will produce and report (not yet available):

- Baseline accuracy of the specialist model before any specialization
- Accuracy after the first specialization cycle, measured against the frozen held-out test set
- Share of volume resolved at acceptable confidence vs. escalated
- Confirmation that 100% of recommendations and promotion decisions are logged and replayable

These numbers will be added to the next revision of this brief once Phase 1 completes.

Summary for Decision Makers

SLM Guard lets institutions run specialist small models on high-volume work with three guarantees that matter in regulated environments:

1. The models stay under strict institutional control — the control plane, not the model, holds decision authority, system access, and the audit record.
2. The models can improve over time without moving sensitive data, using a synthetic-conversion process whose current limits are stated openly, with stronger formal guarantees planned as the system hardens.
3. Every decision and every model change is fully auditable.

This is the operating foundation required to make specialist small models safe, economical, and sustainable at enterprise scale. It is currently at the reference-implementation stage; pilot-scope and engagement discussions are appropriate once Phase 1 results are available.

Contact for discussion of pilot scope or technical deep-dive.